

In [1]: #Import Required Libraries

```
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
```

In [2]: #Load Data

```
df = pd.read_csv("online_retail_II.csv", encoding="ISO-8859-1")
df.head()
```

Out[2]:

	Invoice	StockCode	Description	Quantity	InvoiceDate	Price	Customer ID	Country
0	489434	85048	15CM CHRISTMAS GLASS BALL 20 LIGHTS	12	2009-12-01 07:45:00	6.95	13085.0	United Kingdom
1	489434	79323P	PINK CHERRY LIGHTS	12	2009-12-01 07:45:00	6.75	13085.0	United Kingdom
2	489434	79323W	WHITE CHERRY LIGHTS	12	2009-12-01 07:45:00	6.75	13085.0	United Kingdom
3	489434	22041	RECORD FRAME 7" SINGLE SIZE	48	2009-12-01 07:45:00	2.10	13085.0	United Kingdom
4	489434	21232	STRAWBERRY CERAMIC TRINKET BOX	24	2009-12-01 07:45:00	1.25	13085.0	United Kingdom

In [3]: #Data Cleaning

```
# Remove rows with missing Customer ID
df = df.dropna(subset=['Customer ID'])

# Remove cancelled / invalid transactions
df = df[df['Quantity'] > 0]
df = df[df['Price'] > 0]

# Convert InvoiceDate to datetime
df['InvoiceDate'] = pd.to_datetime(df['InvoiceDate'])

df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
Index: 805549 entries, 0 to 1067370
Data columns (total 8 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Invoice     805549 non-null   object 
 1   StockCode   805549 non-null   object 
 2   Description 805549 non-null   object 
 3   Quantity    805549 non-null   int64  
 4   InvoiceDate 805549 non-null   datetime64[ns]
 5   Price       805549 non-null   float64
 6   Customer ID 805549 non-null   float64
 7   Country     805549 non-null   object 
dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 55.3+ MB
```

```
In [4]: #Feature Engineering - Total Spend
df['Total_Spend'] = df['Quantity'] * df['Price']
```

```
In [5]: #Exploratory Data Analysis (EDA)
#Revenue by Country (Top 10)

country_revenue = (
    df.groupby('Country')['Total_Spend']
    .sum()
    .sort_values(ascending=False)
    .head(10)
)


plt.figure(figsize=(8,5))
country_revenue.plot(kind='bar')
plt.title("Top 10 Countries by Revenue")
plt.ylabel("Revenue")
plt.show()
```



```
In [6]: #Monthly Revenue Trend
df['Month'] = df['InvoiceDate'].dt.to_period('M')


monthly_revenue = df.groupby('Month')['Total_Spend'].sum()

plt.figure(figsize=(10,5))
monthly_revenue.plot()
plt.title("Monthly Revenue Trend")
plt.ylabel("Revenue")
plt.xlabel("Month")
plt.show()
```



```
In [7]: #Top 10 Products by Revenue
top_products = (
    df.groupby('Description')[ 'Total_Spend']
    .sum()
    .sort_values(ascending=False)
    .head(10)
)


plt.figure(figsize=(8,5))
top_products.plot(kind='barh')
plt.title("Top 10 Products by Revenue")
plt.xlabel("Revenue")
plt.show()
```



```
In [8]: #Top Customers by Revenue
```

```
top_customers = (
    df.groupby('Customer ID')['Total_Spend']
    .sum()
    .sort_values(ascending=False)
    .head(10)
)

plt.figure(figsize=(8,5))
top_customers.plot(kind='bar')
plt.title("Top 10 Customers by Revenue")
plt.ylabel("Revenue")
plt.show()
```




In [9]: #Order Value Distribution

```
order_value = df.groupby('Invoice')['Total_Spend'].sum()


plt.figure(figsize=(7,5))
plt.hist(order_value, bins=50)
plt.title("Distribution of Order Value")
plt.xlabel("Order Value")
plt.ylabel("Frequency")
plt.show()
```

### Distribution of Order Value



```
In [10]: #Quantity Distribution  
plt.figure(figsize=(7,4))  
sns.boxplot(x=df['Quantity'])  
plt.title("Quantity Distribution (Outliers)")  
plt.show()
```

### Quantity Distribution (Outliers)



```
In [11]: df.columns
```

```
Out[11]: Index(['Invoice', 'StockCode', 'Description', 'Quantity', 'InvoiceDate',
       'Price', 'Customer ID', 'Country', 'Total_Spend', 'Month'],
      dtype='object')
```

```
In [14]: reference_date = df['InvoiceDate'].max() + pd.Timedelta(days=1)

rfm = df.groupby('Customer ID').agg({
    'InvoiceDate': lambda x: (reference_date - x.max()).days,
    'Invoice': 'nunique',
    'Total_Spend': 'sum'
})

rfm.rename(columns={
    'InvoiceDate': 'Recency',
    'Invoice': 'Frequency',
    'Total_Spend': 'Monetary'
}, inplace=True)

rfm.head()
```

Out[14]:

Customer ID		Recency	Frequency	Monetary
12346.0	326	12	77556.46	
12347.0	2	8	5633.32	
12348.0	75	5	2019.40	
12349.0	19	4	4428.69	
12350.0	310	1	334.40	

```
In [15]: from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
rfm_scaled = scaler.fit_transform(rfm[['Recency', 'Frequency', 'Monetary']])
```

```
In [16]: from sklearn.cluster import MiniBatchKMeans
```

```
kmeans = MiniBatchKMeans(
    n_clusters=4,
    random_state=42,
    batch_size=2048
)

rfm['Cluster'] = kmeans.fit_predict(rfm_scaled)
```

```
In [17]: rfm.head()
```

Out[17]:

**Recency Frequency Monetary Cluster**

<b>Customer ID</b>				
<b>12346.0</b>	326	12	77556.46	2
<b>12347.0</b>	2	8	5633.32	3
<b>12348.0</b>	75	5	2019.40	0
<b>12349.0</b>	19	4	4428.69	3
<b>12350.0</b>	310	1	334.40	1

In [18]:

```
rfm['Customer_Segment'] = rfm['Cluster'].map({
    0: 'High Value Customers',
    1: 'Loyal Customers',
    2: 'At Risk Customers',
    3: 'Low Value Customers'
})
```

In [19]:

```
df = df.merge(
    rfm[['Customer_Segment']],
    left_on='Customer ID',
    right_index=True,
    how='left'
)
```

In [20]:

```
df[['Customer ID', 'Customer_Segment']].head()
```

Out[20]:

**Customer ID Customer\_Segment**

<b>0</b>	13085.0	High Value Customers
<b>1</b>	13085.0	High Value Customers
<b>2</b>	13085.0	High Value Customers
<b>3</b>	13085.0	High Value Customers
<b>4</b>	13085.0	High Value Customers

In [21]:

```
segment_summary = rfm.groupby('Customer_Segment').agg({
    'Recency': 'mean',
    'Frequency': 'mean',
    'Monetary': 'mean',
    'Customer_Segment': 'count'
}).rename(columns={'Customer_Segment': 'Customer_Count'})
```

segment\_summary

Out[21]:


Recency Frequency Monetary Customer\_Count

Customer_Segment		Recency	Frequency	Monetary	Customer_Count
<b>At Risk Customers</b>	26.975535	40.064220	27305.300346		327
<b>High Value Customers</b>	143.022366	3.296645	1160.715084		1699
<b>Low Value Customers</b>	25.671620	7.104012	2707.976407		2019
<b>Loyal Customers</b>	479.968358	2.140753	750.210951		1833

In [22]:

```
segment_revenue = df.groupby('Customer_Segment')['Total_Spend'].sum()

plt.figure(figsize=(6,4))
segment_revenue.plot(kind='bar')
plt.title("Revenue Contribution by Customer Segment")
plt.ylabel("Revenue")
plt.show()
```




In [23]:

```
segment_count = df[['Customer_ID', 'Customer_Segment']].drop_duplicates() \
    .groupby('Customer_Segment').size()

plt.figure(figsize=(6,4))
```

```
segment_count.plot(kind='pie', autopct='%1.1f%%')
plt.title("Customer Distribution by Segment")
plt.ylabel("")
plt.show()
```



In [24]: `!pip install psycopg2-binary sqlalchemy`

```
Requirement already satisfied: psycopg2-binary in c:\users\dell\anaconda3\lib\site-packages (2.9.11)
Requirement already satisfied: sqlalchemy in c:\users\dell\anaconda3\lib\site-packages (2.0.30)
Requirement already satisfied: typing-extensions>=4.6.0 in c:\users\dell\anaconda3\lib\site-packages (from sqlalchemy) (4.15.0)
Requirement already satisfied: greenlet!=0.4.17 in c:\users\dell\anaconda3\lib\site-packages (from sqlalchemy) (3.0.1)
```

In [31]: `from sqlalchemy import create_engine
from sqlalchemy.engine import URL

# Fix Period datatype
df["Month"] = df["Month"].astype(str)

url = URL.create(
 drivername="postgresql+psycopg2",
 username="postgres",
 password="Archana@30",
 host="localhost",
 port=5432,
 database="retail_db"
)

engine = create_engine(url)

engine.connect()
print("Connected successfully!")`

```
df.to_sql(  
    "online_retail_transactions",  
    engine,  
    if_exists="replace",  
    index=False  
)  
  
print("Data successfully loaded into PostgreSQL!")
```

Connected successfully!  
Data successfully loaded into PostgreSQL!

In [ ]: